

# HOUSING PRICE PREDICTION PROJECT



By

Alivia Dasgupta

# Acknowledgement

Firstly, I would like to thank FlipRobo Technologies for giving me the opportunity to work on this project. This project includes mentioning of all the references, research papers, data sources, professionals and other resources that helped me and guided me in completion project.

## INTRODUCTION

### Business Problem Framing

The main objective of this project is to model the price of houses with the available independent variables. This model will then be used by the management to understand how exactly the prices vary with the variables. They can accordingly manipulate the strategy of the firm and concentrate on areas that will yield high returns. Further, the model will be a good way for the management to understand the pricing dynamics of a new market.

### Conceptual Background of the Domain Problem

Houses are one of the necessary need of each and every person around the globe and therefore housing and real estate market is one of the markets which is one of the major contributors in the world's economy. It is a very large market and there are various companies working in the domain. Data science comes as a very important tool to solve problems in the domain to help the companies increase their overall revenue, profits, improving their marketing strategies and focusing on changing trends in house sales and purchases. Predictive modelling, Market mix modelling, recommendation systems are some of the machine learning techniques used for achieving the business goals for housing companies. Our problem is related to one such housing company.

### Technical Requirements

➤ Data contains 1460 entries each having 81 variables.

- Data contains Null values. We need to treat them using the domain knowledge and your own understanding.
- Extensive EDA has to be performed to gain relationships of variable and price.
- Data contains numerical as well as categorical variable. We need to handle them accordingly.
- We have to build Machine Learning methods, apply regularization and determine the optimal values for HyperParameters.
- We need to find important features which affect the price positively or negatively.
- Two datasets are being provided to us (test.csv, train.csv).

## **ADVANTAGES:**

1. The objective behind to take this project is to implement the required data science skills.
2. Improve the analytical thinking.
3. Get into the real world problem solving mechanics.

# **Analytical Problem Framing**

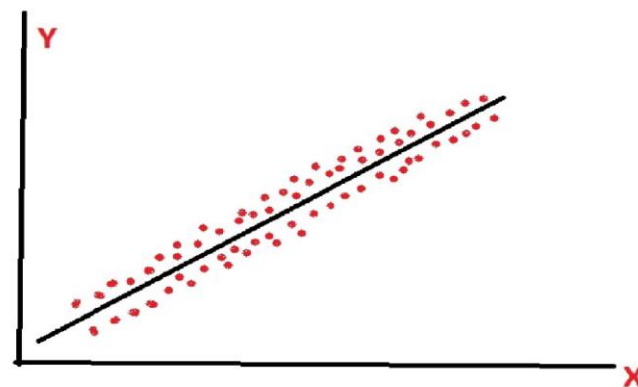
## **Mathematical Modeling of the Problem:**

This is a Regression problem, where our end goal is to predict the Prices of House, based on given data provided in the dataset. We have divided the provided dataset into Training and Testing phases. A Regression Model will be built and trained using the Training data and the Test data is used to predict the outcomes. This will be compared with available test results to find how our model has performed.

We are using Mean Absolute Error, Root Mean Square Error, and 'R2\_Score' to determine the best model among, ❖ Linear Regression ❖ Decision Tree Regression ❖ Random Forest Regression ❖ K Neighbors Regression ❖ Lasso Regression The best results were obtained using Lasso Regression. So, let's discuss a little bit about it.

In a simple regression problem (a single x and a single y), the form of the model would be:

$y = B_0 + B_1 * x$  where  $B_0$  — intercept,  $B_1$  — coefficient,  $x$  — independent variable  $y$  — output or the dependent variable. In higher dimensions when we have more than one input ( $x$ ), The General equation for a Multiple linear regression with  $p$  — independent variables:  $Y = B_0 + B_1 * X_1 + B_2 * X_2 + \dots + B_p * X_p + E(\text{Random Error or Noise})$ .



## Data Sources and their formats

A US-based housing company named Surprise Housing has decided to enter the Australian market. The company uses data analytic to purchase houses at a price below their actual values and flip them at a higher price. For the same purpose, the company has collected a data set from the sale of houses in Australia. The data is provided in the CSV file below. The dataset contains 1460 rows and 81 columns (including the train dataset and test dataset). The top 5 rows of the dataset are:

```
In [5]: df_test.head(10)
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	ScreenPorch	PoolArea	PoolQC	Fence	MiscFeature
0	337	20	RL	86.0	14157	Pave	NaN	IR1	HLS	AllPub	...	0	0	NaN	NaN	Na
1	1018	120	RL	NaN	5814	Pave	NaN	IR1	Lvl	AllPub	...	0	0	NaN	NaN	Na
2	929	20	RL	NaN	11838	Pave	NaN	Reg	Lvl	AllPub	...	0	0	NaN	NaN	Na
3	1148	70	RL	75.0	12000	Pave	NaN	Reg	Bnk	AllPub	...	0	0	NaN	NaN	Na
4	1227	60	RL	86.0	14598	Pave	NaN	IR1	Lvl	AllPub	...	0	0	NaN	NaN	Na
5	650	180	RM	21.0	1936	Pave	NaN	Reg	Lvl	AllPub	...	0	0	NaN	MnPrv	Na
6	1453	180	RM	35.0	3675	Pave	NaN	Reg	Lvl	AllPub	...	0	0	NaN	NaN	Na
7	152	20	RL	107.0	13891	Pave	NaN	Reg	Lvl	AllPub	...	0	0	NaN	NaN	Na
8	427	80	RL	NaN	12800	Pave	NaN	Reg	Low	AllPub	...	396	0	NaN	NaN	Na
9	776	120	RM	32.0	4500	Pave	NaN	Reg	Lvl	AllPub	...	0	0	NaN	NaN	Na

10 rows × 80 columns


```
In [6]: df_train.tail(10)
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal
1158	673	20	RL	NaN	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1159	943	90	RL	42.0	7711	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1160	551	120	RL	53.0	4043	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1161	1301	60	RL	NaN	10762	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1162	1381	30	RL	45.0	8212	Pave	Grvl	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1163	289	20	RL	NaN	9819	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	MnPrv	NaN	0
1164	554	20	RL	67.0	8777	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	MnPrv	NaN	0

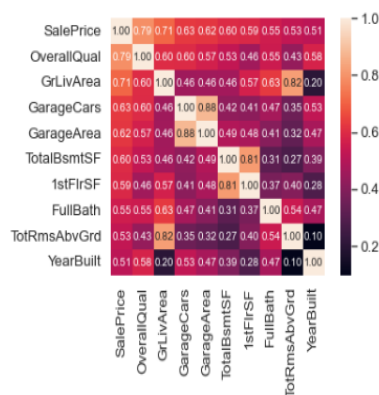
# Data Preprocessing Done

As our dataset contains null values (missing values) so we have replace the missing values with the required values. Details are mentioned below

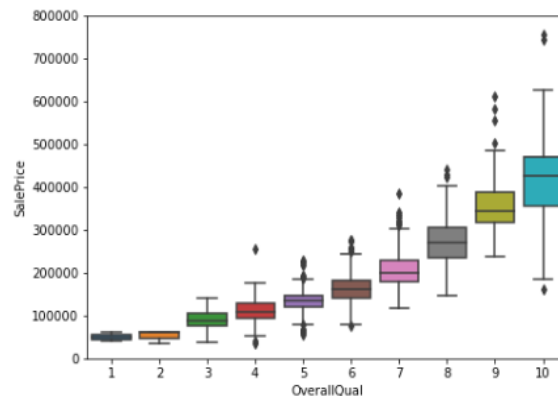
```
In [43]: corrmat = df_train.corr()
f, ax = plt.subplots(figsize=(20, 12))
sns.heatmap(corrmat, vmax=.8, square=True);
```



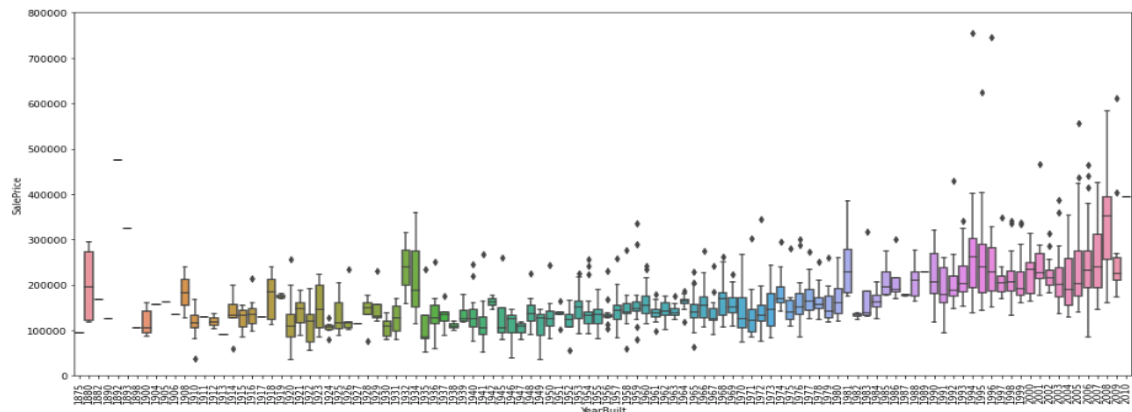
```
In [46]: k = 10 #number of variables for heatmap
cols = corrmat.nlargest(k, 'SalePrice')['SalePrice'].index
cm = np.corrcoef(df_train[cols].values.T)
sns.set(font_scale=1.25)
hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f', annot_kws={'size': 10}, yticklabels=cols.values, xticklabels=
plt.show()
```



```
In [39]: var = 'OverallQual'
data = pd.concat([df_train['SalePrice'], df_train[var]], axis=1)
f, ax = plt.subplots(figsize=(7, 5))
fig = sns.boxplot(x=var, y="SalePrice", data=data)
fig.axis(ymin=0, ymax=800000);
```



```
In [38]: var = 'YearBuilt'
data = pd.concat([df_train['SalePrice'], df_train[var]], axis=1)
f, ax = plt.subplots(figsize=(18, 8))
fig = sns.boxplot(x=var, y="SalePrice", data=data)
fig.axis(ymin=0, ymax=800000);
plt.xticks(rotation=90);
```



## Summary of EDA

1. 'GrLivArea' and 'TotalBsmtSF' seem to be linearly related with 'SalePrice'. Both relationships are positive, which means that as one variable increases, the other also increases. In the case of 'TotalBsmtSF', we can see that the slope of the linear relationship is particularly high.
2. 'OverallQual' and 'YearBuilt' also seem to be related with 'SalePrice'. The relationship seems to be stronger in the case of 'OverallQual', where the box plot shows how sales prices increase with the overall quality. We just analysed four variables, but there are many other that we should analyse. The trick here seems to be the choice of the right features (feature selection) and not the definition of complex relationships between them (feature engineering). That said, let's separate the wheat from the chaff

# Models

```
In [61]: from sklearn.preprocessing import LabelEncoder  
lblencoder = LabelEncoder()
```

```
In [62]: trn_data = df_train.copy()  
tst_data = df_test.copy()
```

```
In [63]: for col in trn_data.select_dtypes(include=['O']).columns:  
    trn_data[col] = lblencoder.fit_transform(trn_data[col].astype(str))  
    tst_data[col] = lblencoder.fit_transform(tst_data[col].astype(str))
```

```
In [64]: targets = trn_data['SalePrice']  
trn_data.drop(['SalePrice'], axis=1, inplace=True)  
  
from sklearn.preprocessing import StandardScaler, MinMaxScaler  
mmmx = MinMaxScaler()  
stdnorm = StandardScaler()
```

```
In [65]: def standardized_normalized(data):  
    df_col = data.columns  
    data = mmmx.fit_transform(data)  
    data = stdnorm.fit_transform(data)  
    return pd.DataFrame(data, columns=df_col)
```

```
In [66]: trn_data.shape
```

```
Out[66]: (1168, 75)
```

```
In [67]: trn_data = standardized_normalized(trn_data)  
tst_data = standardized_normalized(tst_data)
```

```
In [68]: trn_data.sample(10)
```

```
Out[68]:
```

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	LandContour	Utilities	LotConfig	LandSlope	...	OpenPorchSF	EnclosedPorch	3S
537	2.939506	1.558511	-1.934561	-0.999793	0.058621	0.752055	0.318473	0.0	0.606420	-0.226126	...	-0.701705	-0.364375	-1
1053	-0.877042	-0.021646	10.307645	1.917132	0.058621	-0.664720	-2.499530	0.0	0.606420	3.295414	...	-0.701705	-0.364375	-1
674	0.077095	-0.021646	0.539035	-0.027335	0.058621	0.752055	0.318473	0.0	0.606420	-0.226126	...	0.081987	-0.364375	-1
237	-0.877042	-0.021646	-0.718726	0.601462	0.058621	-1.373107	0.318473	0.0	-1.829688	-0.226126	...	-0.219433	-0.364375	-1
421	1.508301	-0.021646	-1.389531	-0.605538	0.058621	0.752055	0.318473	0.0	0.606420	-0.226126	...	0.232697	-0.364375	-1
886	-0.877042	-0.021646	-1.012203	0.182413	0.058621	-1.373107	0.318473	0.0	-1.220661	-0.226126	...	0.021703	-0.364375	-1
459	2.462438	-0.021646	-1.808785	-0.913235	0.058621	0.752055	0.318473	0.0	0.606420	-0.226126	...	-0.701705	-0.364375	-1
462	-0.877042	-3.181960	-0.718726	-0.679251	0.058621	-0.664720	0.318473	0.0	0.606420	-0.226126	...	1.468519	-0.364375	-1
24	-0.877042	-0.021646	0.161707	-0.140139	0.058621	0.752055	0.318473	0.0	-1.829688	-0.226126	...	0.262839	-0.364375	-1
189	-0.877042	-0.021646	-0.718726	0.224295	0.058621	-1.373107	0.318473	0.0	0.606420	-0.226126	...	-0.701705	-0.364375	-1

```
In [72]: from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
from xgboost import XGBRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import cross_validate, ShuffleSplit
```

```
In [73]: ml_algo = [
    LinearRegression(),
    DecisionTreeRegressor(),
    RandomForestRegressor(),
    SVR(),
    XGBRegressor(),
    KNeighborsRegressor()
]
```

```
In [74]: ml_columns = ['MLAlgo Name', 'MLAlgo Parameters', 'MLAlgo Train MAE', 'MLAlgo Test MAE', 'MLAlgo Time']
ml_algo_compare = pd.DataFrame(columns = ml_columns)

ml_algo_predict = pd.DataFrame(targets.copy().values, columns=['Actual Sales'])
```

```
In [75]: row_index=0
for alg in ml_algo:
    #set name and parameters
    MLA_name = alg.__class__.__name__
    ml_algo_compare.loc[row_index, 'MLAlgo Name'] = MLA_name
    ml_algo_compare.loc[row_index, 'MLAlgo Parameters'] = str(alg.get_params())
    cv_results = cross_validate(alg, trn_data, targets, scoring='neg_mean_absolute_error', cv = 3, return_train_score=True)
    ml_algo_compare.loc[row_index, 'MLAlgo Time'] = cv_results['fit_time'].mean()
    ml_algo_compare.loc[row_index, 'MLAlgo Train MAE'] = cv_results['train_score'].mean()
    ml_algo_compare.loc[row_index, 'MLAlgo Test MAE'] = cv_results['test_score'].mean()
    alg.fit(trn_data, targets)
    ml_algo_predict[MLA_name] = alg.predict(trn_data)
    row_index += 1
```

```
In [78]: dtmodel = XGBRegressor()
dtmodel.fit(trn_data, targets)
```

```
Out[78]: XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None,
    colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
    early_stopping_rounds=None, enable_categorical=False,
    eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',
    importance_type=None, interaction_constraints='',
    learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,
    max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,
    missing=nan, monotone_constraints='()', n_estimators=100, n_jobs=0,
    num_parallel_tree=1, predictor='auto', random_state=0, reg_alpha=0,
    reg_lambda=1, ...)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [79]: from sklearn.metrics import mean_absolute_error, r2_score
trn_p = dtmodel.predict(trn_data)
tst_ = dtmodel.predict(tst_data)
```

```
In [80]: print("training mae:", mean_absolute_error(targets, trn_p))
print("training r2:", r2_score(targets, trn_p))
```

```
training mae: 695.7653306934932
training r2: 0.9998312585722334
```

In [ ]:



## **Conclusion**

We have observed about the variables, we analysed 'SalePrice' alone and with the most correlated variables, we dealt with missing data and outliers, we tested some of the fundamental statistical assumptions and we even transformed categorial variables into dummy variables. That's a lot of work that Python helped us make easier.